

¿Quiénes somos?

Doctor en Ciencia y Tecnología Informática

Inteligencia artificial, Robótica, Espacio, Videojuegos y Aprendizaje Automático

@moisipm

Moisés Martínez



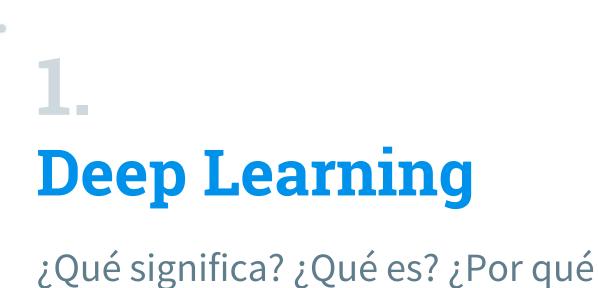
Estudiante de Doctorado en Ciencia y Tecnología Informática

Inteligencia artificial, Robótica y Planificación Automática Multi Agente

@sailormercury91

Nerea Luis





todos hablan de deep learning?



¿Qué significa Machine Learning?

Conjunto de modelos y algoritmos basados en funciones matemáticas que nos permiten identificar patrones con el fin de predecir

Tipos de aprendizaje



Aprendizaje supervisado con intervención humana

¿Cómo clasifico fotos de gatos y perros?

Aprendizaje no supervisado

sin intervención humana

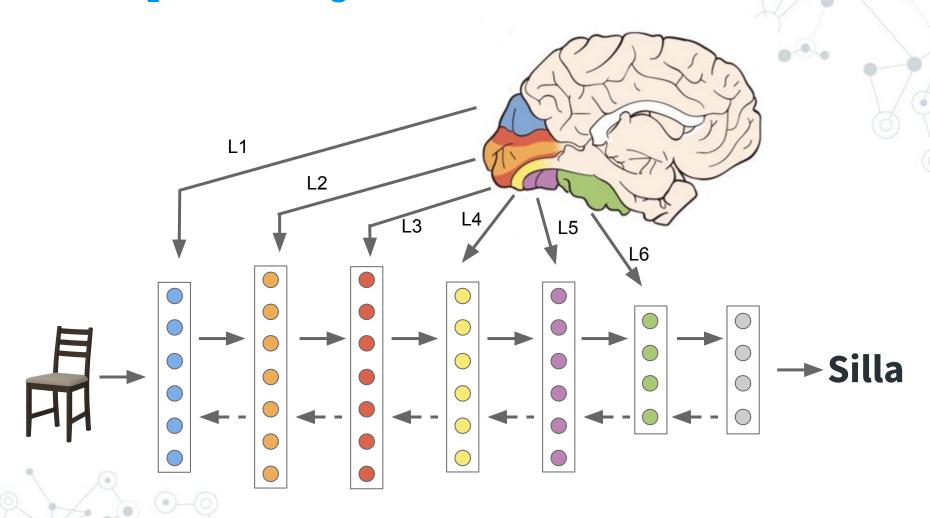
¿Cómo clasifico tipos de perros si no sé qué los distinguen?

Aprendizaje por refuerzo sin intervención humana

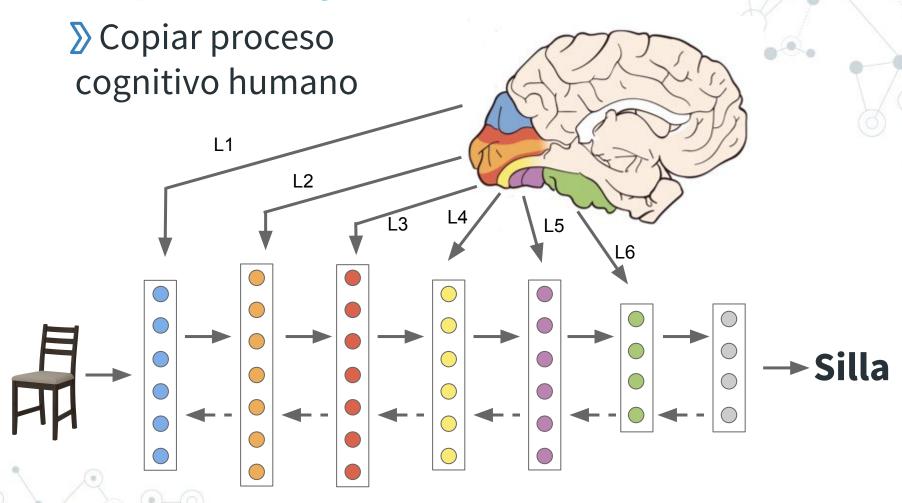
¿Cómo entreno a mi perro? ¿Robótica?



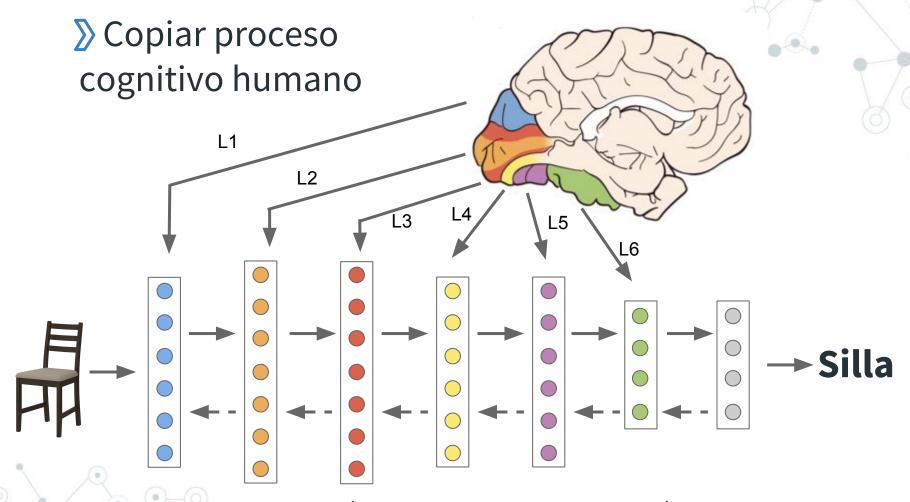
Deep Learning



Deep Learning



Deep Learning



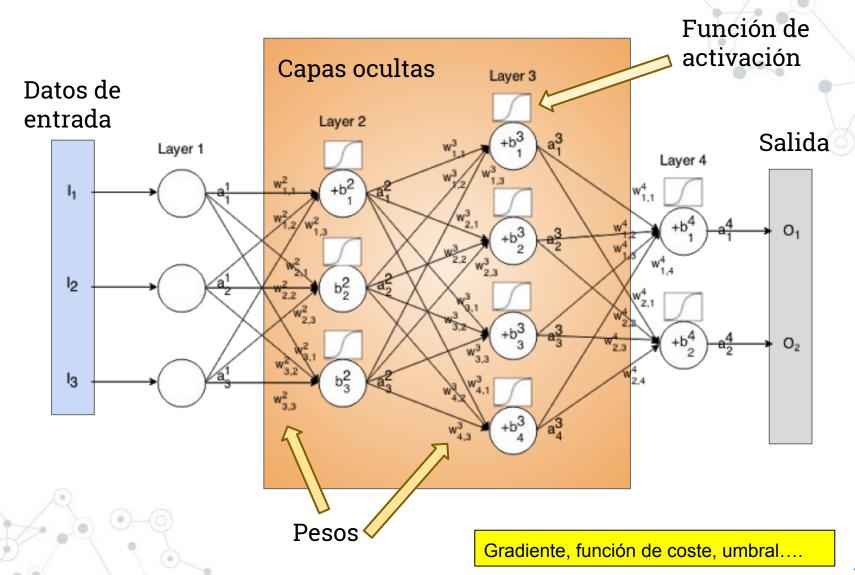
Colección de funciones matemáticas sencillas

Entendiendo una red de neuronas

- > Utilizamos elementos llamados 'neuronas'
- > Conectadas de forma múltiple entre sí formando una 'red'
- > Damos importancia a los 'datos de entrada'
- > Escogemos pesos, función de activación
- > Interpretamos la salida:
 - Semántica de los datos
 - Características Comunes



¡Y ahora todo junto que empezamos!

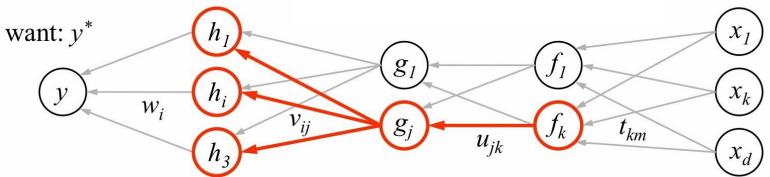


Ejemplos de funciones de activación

ReLU

Name	Plot	Equation
Identity	/	f(x) = x
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \ge 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \ge 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) ^[2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \ge 0 \end{cases}$
Exponential Linear Unit (ELU) ^[3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \ge 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$

El algoritmo de back-propagation



- 1. receive new observation $\mathbf{x} = [x_1 ... x_d]$ and target \mathbf{y}^*
- 2. **feed forward:** for each unit g_j in each layer 1...L compute g_j based on units f_k from previous layer: $g_j = \sigma \left(u_{j0} + \sum_k u_{jk} f_k \right)$
- 3. get prediction y and error $(y-y^*)$
- **4.** back-propagate error: for each unit g_i in each layer L...1
- (a) compute error on g_j

$$\frac{\partial E}{\partial g_{j}} = \sum_{i} \sigma'(h_{i}) v_{ij} \frac{\partial E}{\partial h_{i}}$$
should g_{j} how h_{i} will was h_{i} too be higher change as high or or lower? g_{i} changes too low?

- (b) for each u_{jk} that affects g_j
 - (i) compute error on u_{jk}

$$\frac{\partial E}{\partial u_{jk}} = \frac{\partial E}{\partial g_{j}} \sigma'(g_{j}) f_{k}$$

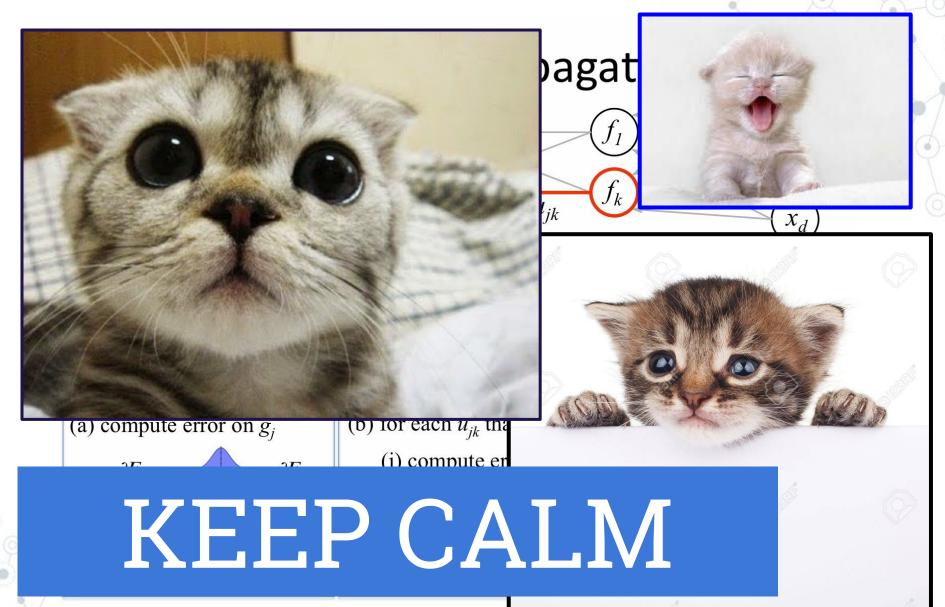
do we want g_j to he be higher/lower if

 $f_{k} \qquad u_{jk} \leftarrow u_{jk} - \eta \frac{\partial E}{\partial u_{jk}}$

how g_j will change if u_{jk} is higher/lower

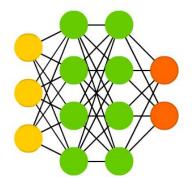
(ii) update the weight

WTF IS THIS SH*T PLEASE STOP NOW

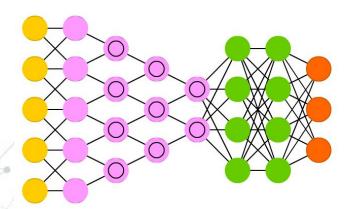


Tipos de redes de neuronas... ...para trabajar en <u>Deep</u> Learning

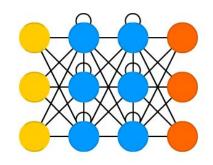
Feed Forward (FFN)



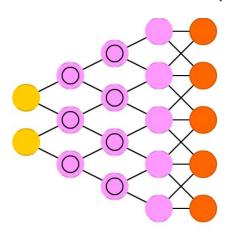
Convolucionales (CNN)



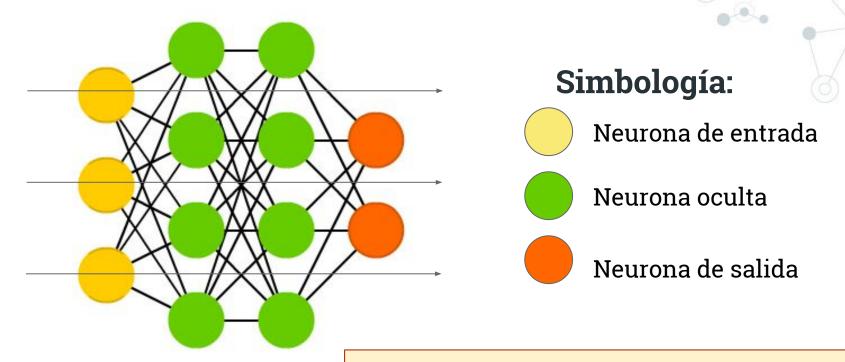
Recurrentes (RNN)



Deconvolucionales (DNN)

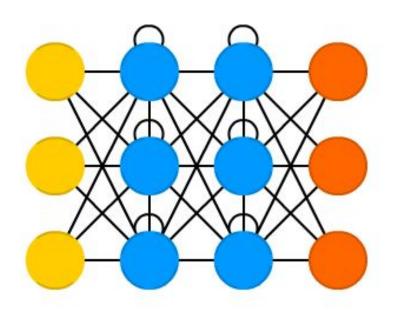


Feed Forward Neural Networks (FFNN)

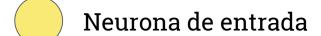


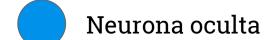
Todo el flujo de operaciones se realiza desde los datos de entrada hacia la salida en la misma dirección. De izquierda a derecha, 'avanzando hacia adelante'

Recurrent Neural Networks (RNN)



Simbología:





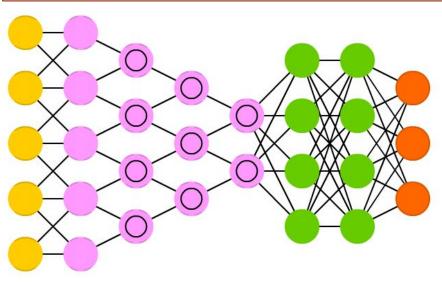


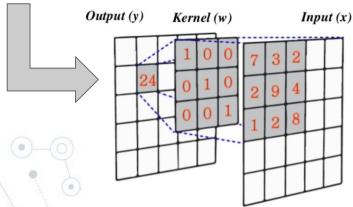
Las neuronas de las capas ocultas retienen más tiempo el entrenamiento si el peso es elevado. Esto es porque el propio output se recibe de nuevo como input.

Convolutional Neural Networks (CNN)

Matriz de mxn pesos. Equivalentes cada uno a un trozo pequeño de la imagen

I M A G E N





Simbología:

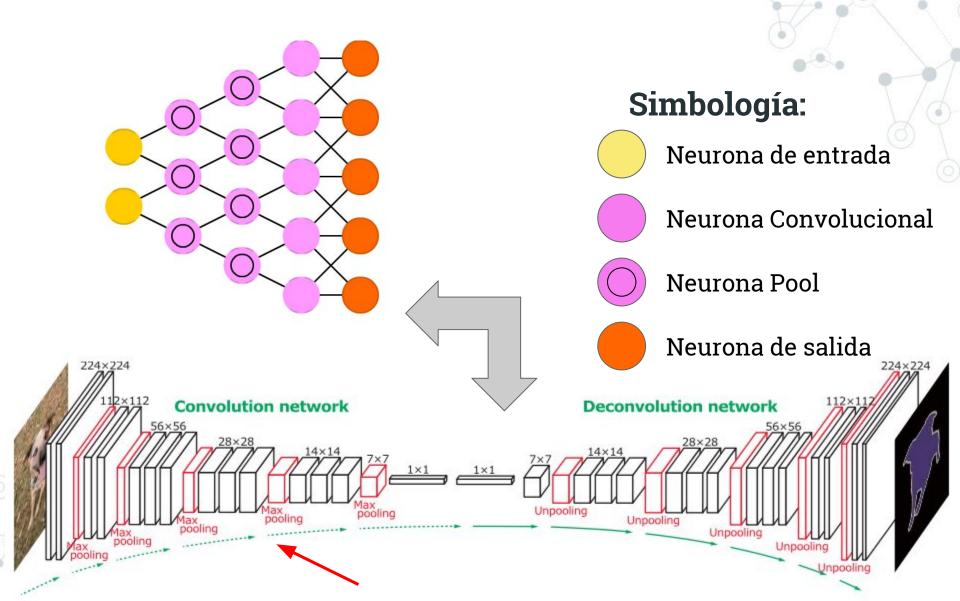
- Neurona de entrada
- Neurona Convolucional
- Neurona Pool
- Neurona Forward
- Neurona de salida

Permite submuestreo de datos de entrada.

Convolutional Neural Networks (CNN)

Matriz de mxn pesos. Equivalentes Simbología: cada uno a un trozo pequeño de la imagen Neurona de entrada Neurona Convolucional M A Neurona Pool G E Neurona Forward N Neurona de salida 19 Capa 1 Capa 2 Capa 3 Capa 4

Deconvolutional Neural Networks (DNN)



WTF IS THIS SH*T AGAIN PLEASE STOP





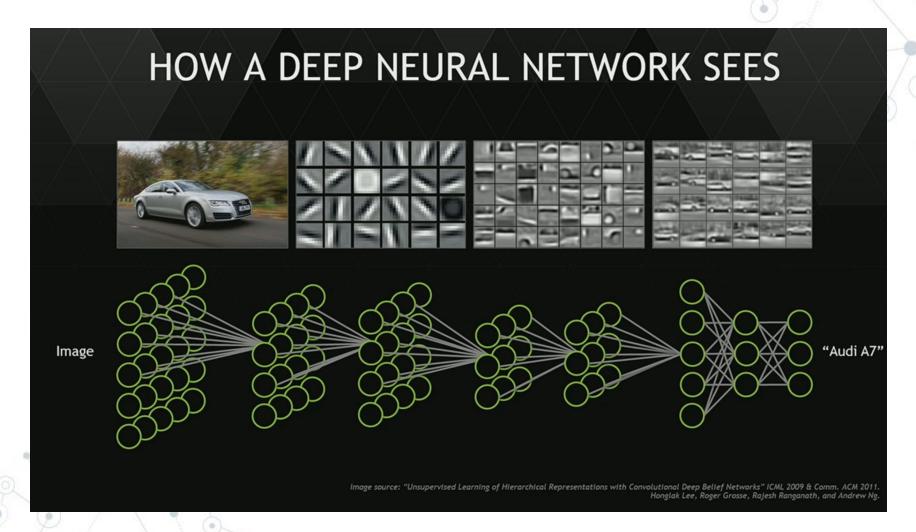


KEEP CALM

¿Pero entonces qué es <u>Deep</u> Learning?

hidden layer 1 hidden layer 2 hidden layer 3 input layer output layer

¿Pero entonces qué es <u>Deep</u> Learning?



Un par de 'detallitos' o cinco....

- Las redes de neuronas se encuentran dentro de las técnicas de Inteligencia Artificial de 'caja negra'
- Hay que probar y probar arquitecturas de redes
- No es evidente por qué unas funcionan mejor y otras no
- No toda tu solución tiene que estar centrada en DL o NN hay más técnicas
- DL no es aprendizaje no supervisado y viceversa



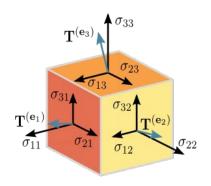
2.

TensorFlow

Abramos la caja negra de la que todo el mundo habla

Tensor \longrightarrow **N-dimensional array**

Vector 1-D tensor

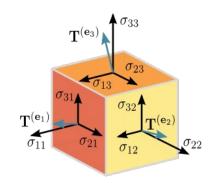




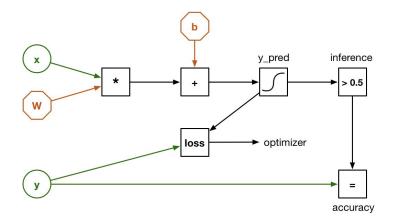
Tensor \longrightarrow **N-dimensional array**

Vector 1-D tensor

Matriz \Longrightarrow 2-D tensor



Flow Secuencia de operaciones



Rectified Linear Unit (ReLU)
$$f(x) = \max(0,x)$$

Funcion de activacion

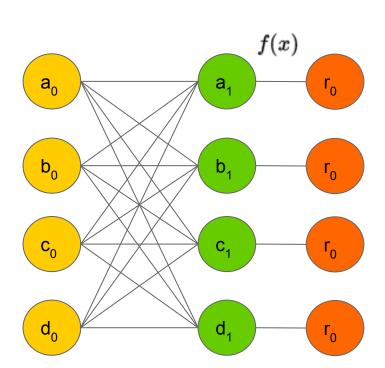


Rectified Linear Unit (ReLU)



$$f(x) = \max(0, x)$$

Funcion de activacion

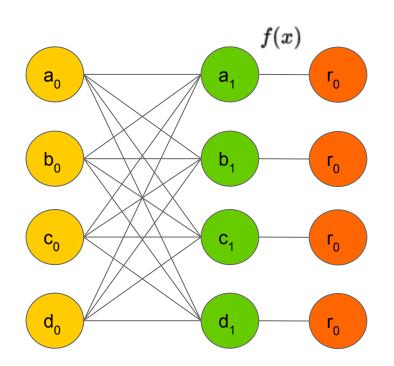


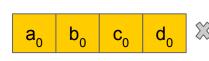
Rectified Linear Unit (ReLU)



$$f(x) = \max(0, x)$$

Funcion de activacion





f ₀₀	f ₀₁	f ₀₂	f ₀₃
f ₁₀	f ₁₁	f ₁₂	f ₁₃
f ₂₀	f ₂₁	f ₂₂	f ₂₃
f ₃₀	f ₃₁	f ₂₃	f ₃₃

$$a_1 = a_0 f_{00} + b_0 f_{10} + c_0 f_{20} + d_0 f_{30}$$

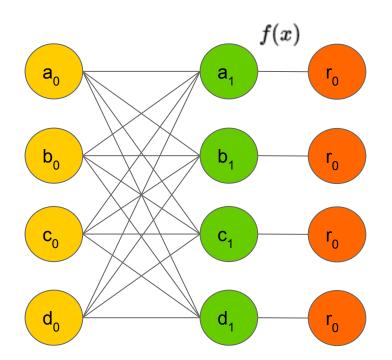
$$d_1 = a_0 f_{03} + b_0 f_{13} + c_0 f_{23} + d_0 f_{33}$$

Rectified Linear Unit (ReLU)



$$f(x) = \max(0, x)$$

Funcion de activacion



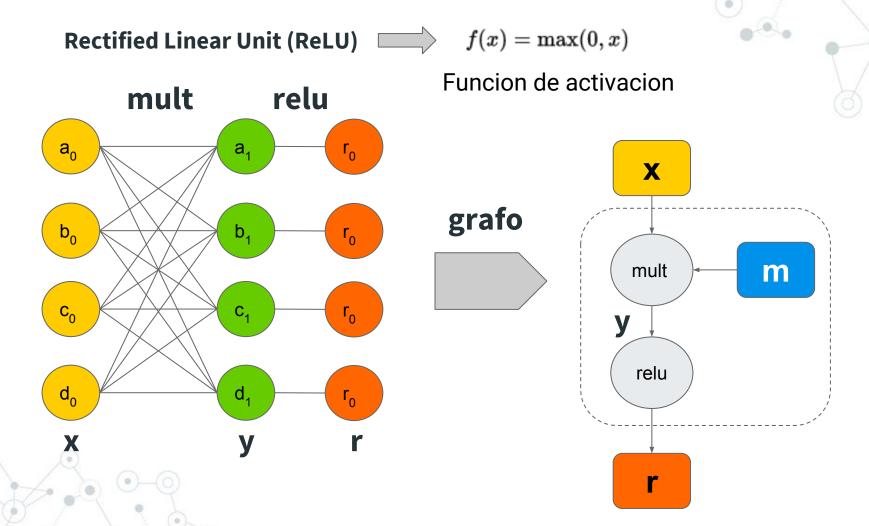
ReLU $\begin{bmatrix} a_1 & b_1 & c_1 & d_1 \end{bmatrix}$

$$r_1 = \max(0, a_1)$$

•

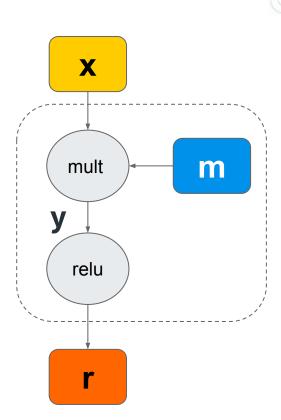
.

 $r_1 = \max(0, a_1)$



Rectified Linear Unit (ReLU) $f(x) = \max(0,x)$

```
import tensorflow as tf
import numpy as np
session = tf.Session()
x = tf.placeholder("float", [1,4])
m = tf.Variable(tf.random_normal([4,4]), name='m')
y = tf.matmul(x, m)
r = tf.nn.relu(y)
session.run(tf.global_variables_initializer())
print(session.run(r, feed_dict {x:np.array([[1.0,2.0,3.0,4.0]])}))
session.close()
```



Rectified Linear Unit (ReLU) $f(x) = \max(0,x)$

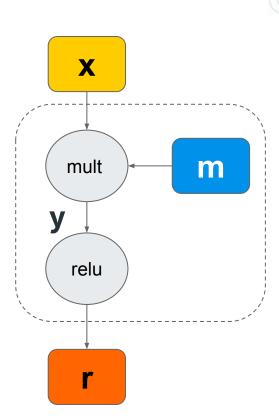
import tensorflow as tf
import numpy as np

session = tf.Session()

x = tf.placeholder("float", [1,4])
m = tf.Variable(tf.random_normal([4,4]), name='m')
y = tf.matmul(x, m)
r = tf.nn.relu(y)

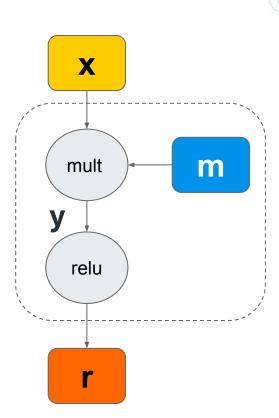
session.run(tf.global_variables_initializer())

print(session.run(r, feed_dict {x:np.array([[1.0,2.0,3.0,4.0]])})))
session.close()



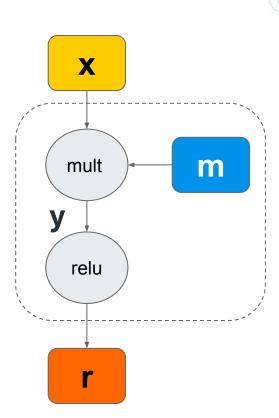
Rectified Linear Unit (ReLU)
$$f(x) = \max(0,x)$$

```
import tensorflow as tf
import numpy as np
session = tf.Session()
x = tf.placeholder("float", [1,4])
m = tf.Variable(tf.random_normal([4,4]), name='m')
y = tf.matmul(x, m)
r = tf.nn.relu(y)
session.run(tf.global_variables_initializer())
print(session.run(r, feed_dict {x:np.array([[1.0,2.0,3.0,4.0]])}))
session.close()
```



Rectified Linear Unit (ReLU) $f(x) = \max(0,x)$

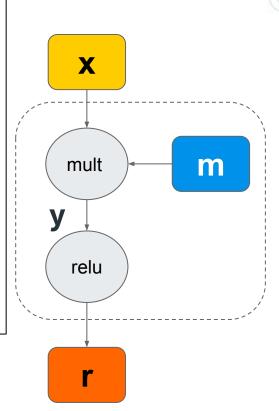
```
import tensorflow as tf
import numpy as np
session = tf.Session()
x = tf.placeholder("float", [1,4])
m = tf.Variable(tf.random_normal([4,4]), name='m')
y = tf.matmul(x, m)
r = tf.nn.relu(y)
session.run(tf.global_variables_initializer())
print(session.run(r, feed_dict {x:np.array([[1.0,2.0,3.0,4.0]])}))
session.close()
```



Tensor Flow

Rectified Linear Unit (ReLU) $f(x) = \max(0,x)$

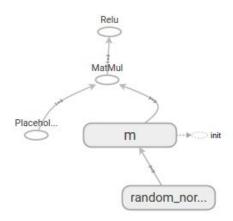
import tensorflow as tf import numpy as np session = tf.Session() x = tf.placeholder("float", [1,4]) m = tf.Variable(tf.random_normal([4,4]), name='m') y = tf.matmul(x, m)r = tf.nn.relu(y)session.run(tf.global_variables_initializer()) print(session.run(r, feed_dict {x:np.array([[1.0,2.0,3.0,4.0]])})) session.close()



Tensor Flow

Rectified Linear Unit (ReLU)
$$f(x) = \max(0,x)$$

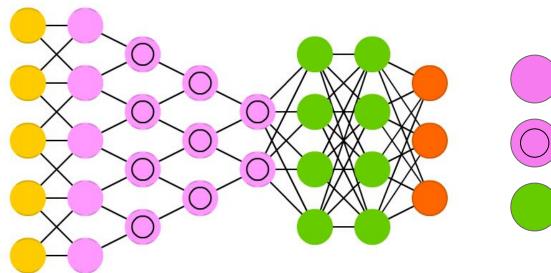
```
import tensorflow as tf
import numpy as np
session = tf.Session()
x = tf.placeholder("float", [1,4])
m = tf.Variable(tf.random_normal([4,4]), name='m')
y = tf.matmul(x, m)
r = tf.nn.relu(y)
session.run(tf.global_variables_initializer())
print(session.run(r, feed_dict {x:np.array([[1.0,2.0,3.0,4.0]])}))
session.close()
```



3. Creando mi Red de Neuronas

¿Cómo la creo? ¿Qué necesito?

Redes Convolucionales



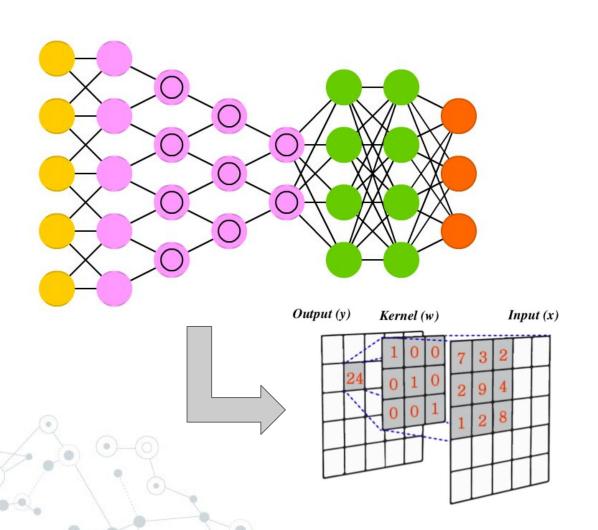
Operaciones



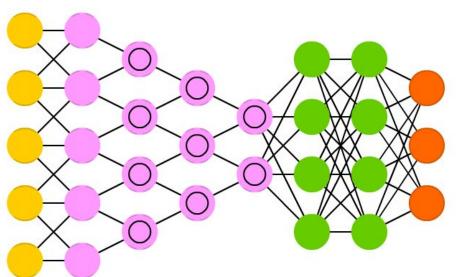




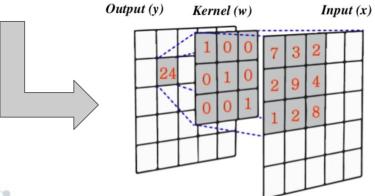
Redes Convolucionales: Convoluciones



Redes Convolucionales: Convoluciones

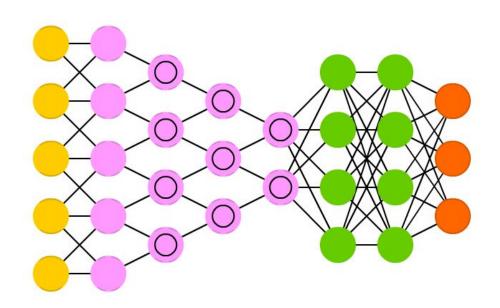


conv = tf.layers.conv2d(
 inputs=input_layer,
 filters=64,
 w=[6, 6],
 padding="same",
 activation=tf.nn.relu)

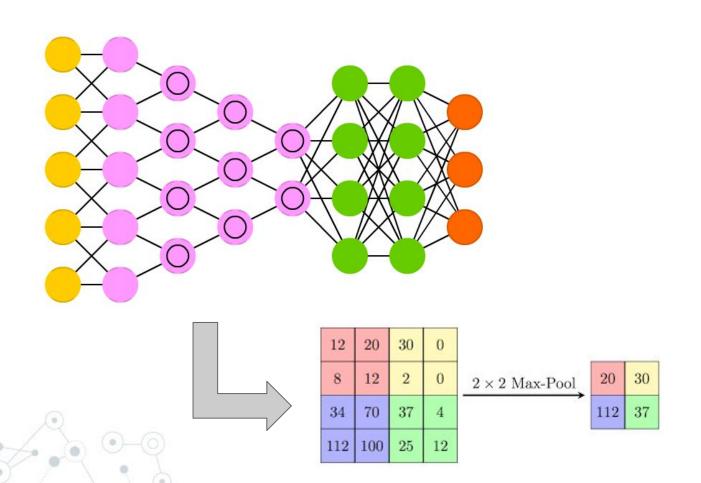




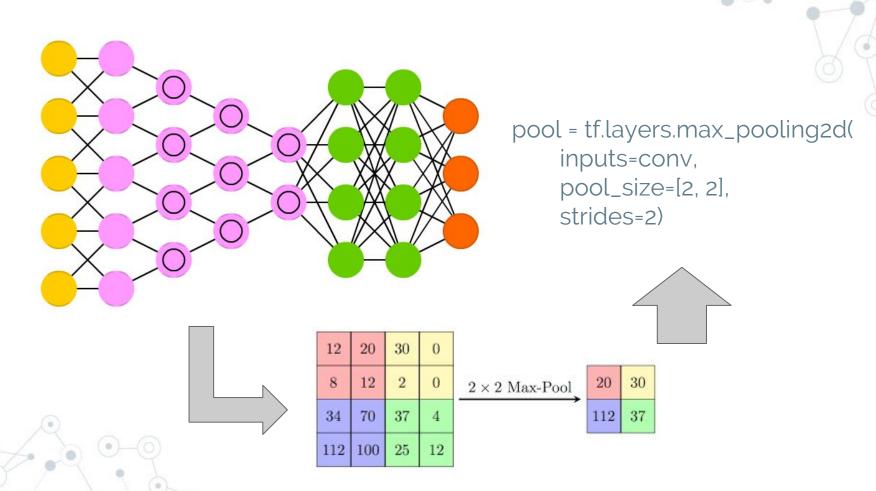
Redes Convolucionales: Max - Pooling



Redes Convolucionales: Max - Pool



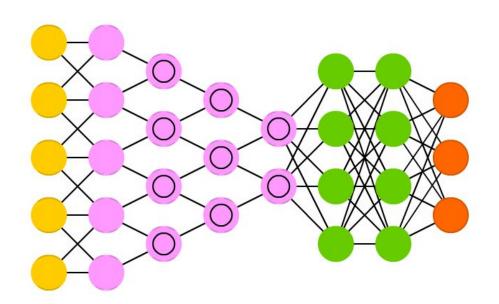
Redes Convolucionales: Max - Pooling



Redes Convolucionales: Dense

Extraer información

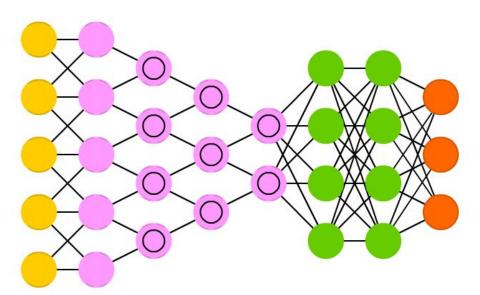
Clasificar información



dense = tf.layers.dense(
 inputs=pool
 units=1024,
 activation=tf.nn.relu)

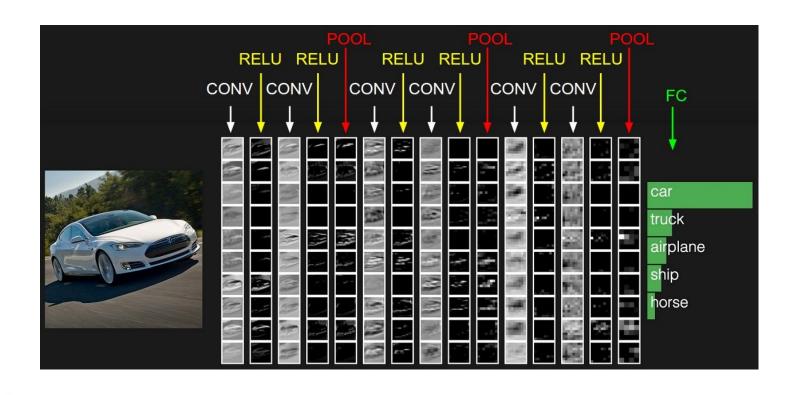
Redes Convolucionales: Dense

Extraer información Clasificar información



pool_reshape = tf.reshape(pool, [-1, 7 * 7 * 64])
dense = tf.layers.dense(
 inputs=pool_reshape
 units=1024,
 activation=tf.nn.relu)

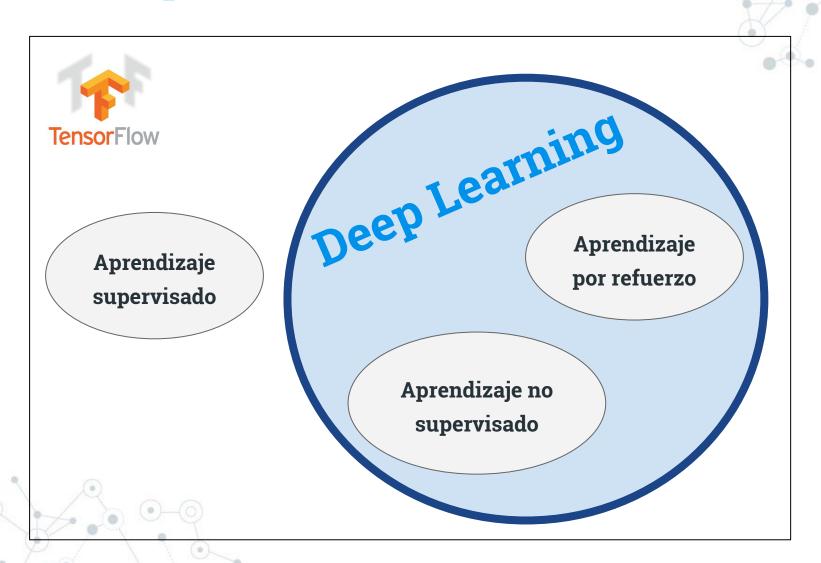
Redes Convolucionales



4. Son sólo redes de neuronas

¿Podemos construir otros algoritmos?

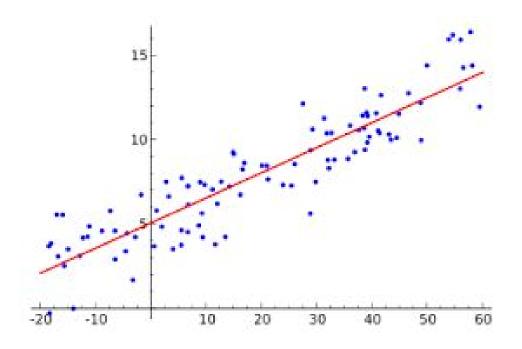
Más que redes de neuronas



Regresión Lineal

Construye un modelo para aproximar la relación de dependencia entre una variable dependiente Y y conjunto de variables independientes X.

$$Y_t = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \varepsilon$$

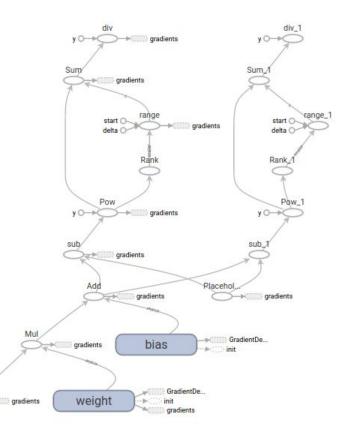


```
import tensorflow as tf
                                                  Y = w * X + b
import numpy
rng = numpy.random
X = tf.placeholder("float")
Y = tf.placeholder("float")
w = tf.Variable(rng.randn(), name="weight")
b = tf.Variable(rng.randn(), name="bias")
operation = tf.add(tf.multiply(X, w), b)
cost = tf.reduce sum(tf.pow(operation-Y, 2))/(2*n samples)
                                                                                                                                       Rank/
optimizer = tf.train.GradientDescentOptimizer(learning rate).minimize(cost)
with tf.Session() as session:
                                                                                                                gradients
  session.run(tf.global variables initializer())
  for epoch in range(training epochs):
     for (x, y) in zip(train X, train Y):
                                                                                                         gradients
       sess.run(optimizer, feed_dict={X: x, Y: y})
     if (epoch+1) % display step == 0:
       c = session.run(cost, feed dict={X: train X, Y:train Y})
                                                                                                             gradients
  training cost = session.run(cost, feed dict={X: train X, Y: train Y})
  testing cost = session.run(
                                                                                                               bias
                                                                                                 gradients
     tf.reduce sum(tf.pow(operation-Y, 2)) / (2 * test X.shape[0]),
     feed dict={X: test X, Y: test Y})
                                                                                                                   GradientDe.
                                                                                                  weight
                                                                                     gradients
```

import numpy rng = numpy.random X = tf.placeholder("float") Variables de entrada Y = tf.placeholder("float") w = tf.Variable(rng.randn(), name="weight") b = tf.Variable(rng.randn(), name="bias") operation = tf.add(tf.multiply(X, w), b) cost = tf.reduce sum(tf.pow(operation-Y, 2))/(2*n samples) optimizer = tf.train.GradientDescentOptimizer(learning rate).minimize(cost) with tf.Session() as session: session.run(tf.global variables initializer()) for epoch in range(training epochs): for (x, y) in zip(train X, train Y): sess.run(optimizer, feed_dict={X: x, Y: y}) if (epoch+1) % display step == 0: c = session.run(cost, feed_dict={X: train_X, Y:train_Y}) training cost = session.run(cost, feed dict={X: train X, Y: train Y}) testing cost = session.run(tf.reduce sum(tf.pow(operation-Y, 2)) / (2 * test X.shape[0]),

feed dict={X: test X, Y: test Y})

import tensorflow as tf



```
import tensorflow as tf
import numpy
rng = numpy.random
X = tf.placeholder("float")
Y = tf.placeholder("float")
w = tf.Variable(rng.randn(), name="weight")
b = tf.Variable(rng.randn(), name="bias")
operation = tf.add(tf.multiply(X, w), b) Recta de regresión
cost = tf.reduce_sum(tf.pow(operation-Y, 2))/(2*n_samples)
optimizer = tf.train.GradientDescentOptimizer(learning rate).minimize(cost)
                                                                                                                                        Rank/
with tf.Session() as session:
  session.run(tf.global variables initializer())
                                                                                                                 gradients
  for epoch in range(training epochs):
     for (x, y) in zip(train X, train Y):
       sess.run(optimizer, feed_dict={X: x, Y: y})
                                                                                                          gradients
     if (epoch+1) % display step == 0:
       c = session.run(cost, feed dict={X: train X, Y:train Y})
  training cost = session.run(cost, feed dict={X: train X, Y: train Y})
                                                                                                              gradients
  testing cost = session.run(
     tf.reduce sum(tf.pow(operation-Y, 2)) / (2 * test X.shape[0]),
                                                                                                                bias
                                                                                                  gradients
     feed dict={X: test X, Y: test Y})
                                                                                                                    GradientDe.
                                                                                      gradients
                                                                                                  weight
```

```
import tensorflow as tf
import numpy
rng = numpy.random
X = tf.placeholder("float")
Y = tf.placeholder("float")
w = tf.Variable(rng.randn(), name="weight")
b = tf.Variable(rng.randn(), name="bias")
operation = tf.add(tf.multiply(X, w), b)
cost = tf.reduce sum(tf.pow(operation-Y, 2))/(2*n samples)
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
                                                                                                                                         Rank/
                                              Algoritmo
with tf.Session() as session:
  session.run(tf.global variables initializer())
                                                                                                                 gradients
  for epoch in range(training epochs):
     for (x, y) in zip(train X, train Y):
        sess.run(optimizer, feed_dict={X: x, Y: y})
                                                                                                           gradients
     if (epoch+1) % display step == 0:
        c = session.run(cost, feed dict={X: train_X, Y:train_Y})
  training cost = session.run(cost, feed dict={X: train X, Y: train Y})
                                                                                                               gradients
  testing cost = session.run(
     tf.reduce sum(tf.pow(operation-Y, 2)) / (2 * test X.shape[0]),
                                                                                                                 bias
                                                                                                  gradients
     feed dict={X: test X, Y: test Y})
                                                                                                                     GradientDe.
                                                                                      gradients
                                                                                                   weight
```

Conclusiones ¿Qué has aprendido en 35 minutos?

Conclusiones

- TensorFlow nos permite construir diferentes tipos de algoritmos
- Las redes convolucionales son complicadas
- No es sólo sirve para crear redes de neuronas.
- El algoritmo se transforma en un grafo de operaciones dependiendo de los datos (tensores)
- Las operaciones sobre los tensores se paralelizan

http://t3chfest.uc3m.es

4 TRACKS | 90 CHARLAS/TALLERES | CONCURSO PROGRAMACIÓN | HACKATHON | GOOGLE HASH CODE | 1500 PERSONAS... ¡Y FALTAS







¡Envíanos tu propuesta! DEADLINE: DICIEMBRE

¡Muchas Gracias!

Preguntas?



@moisipm



@sailormercury91
@witsstories

